

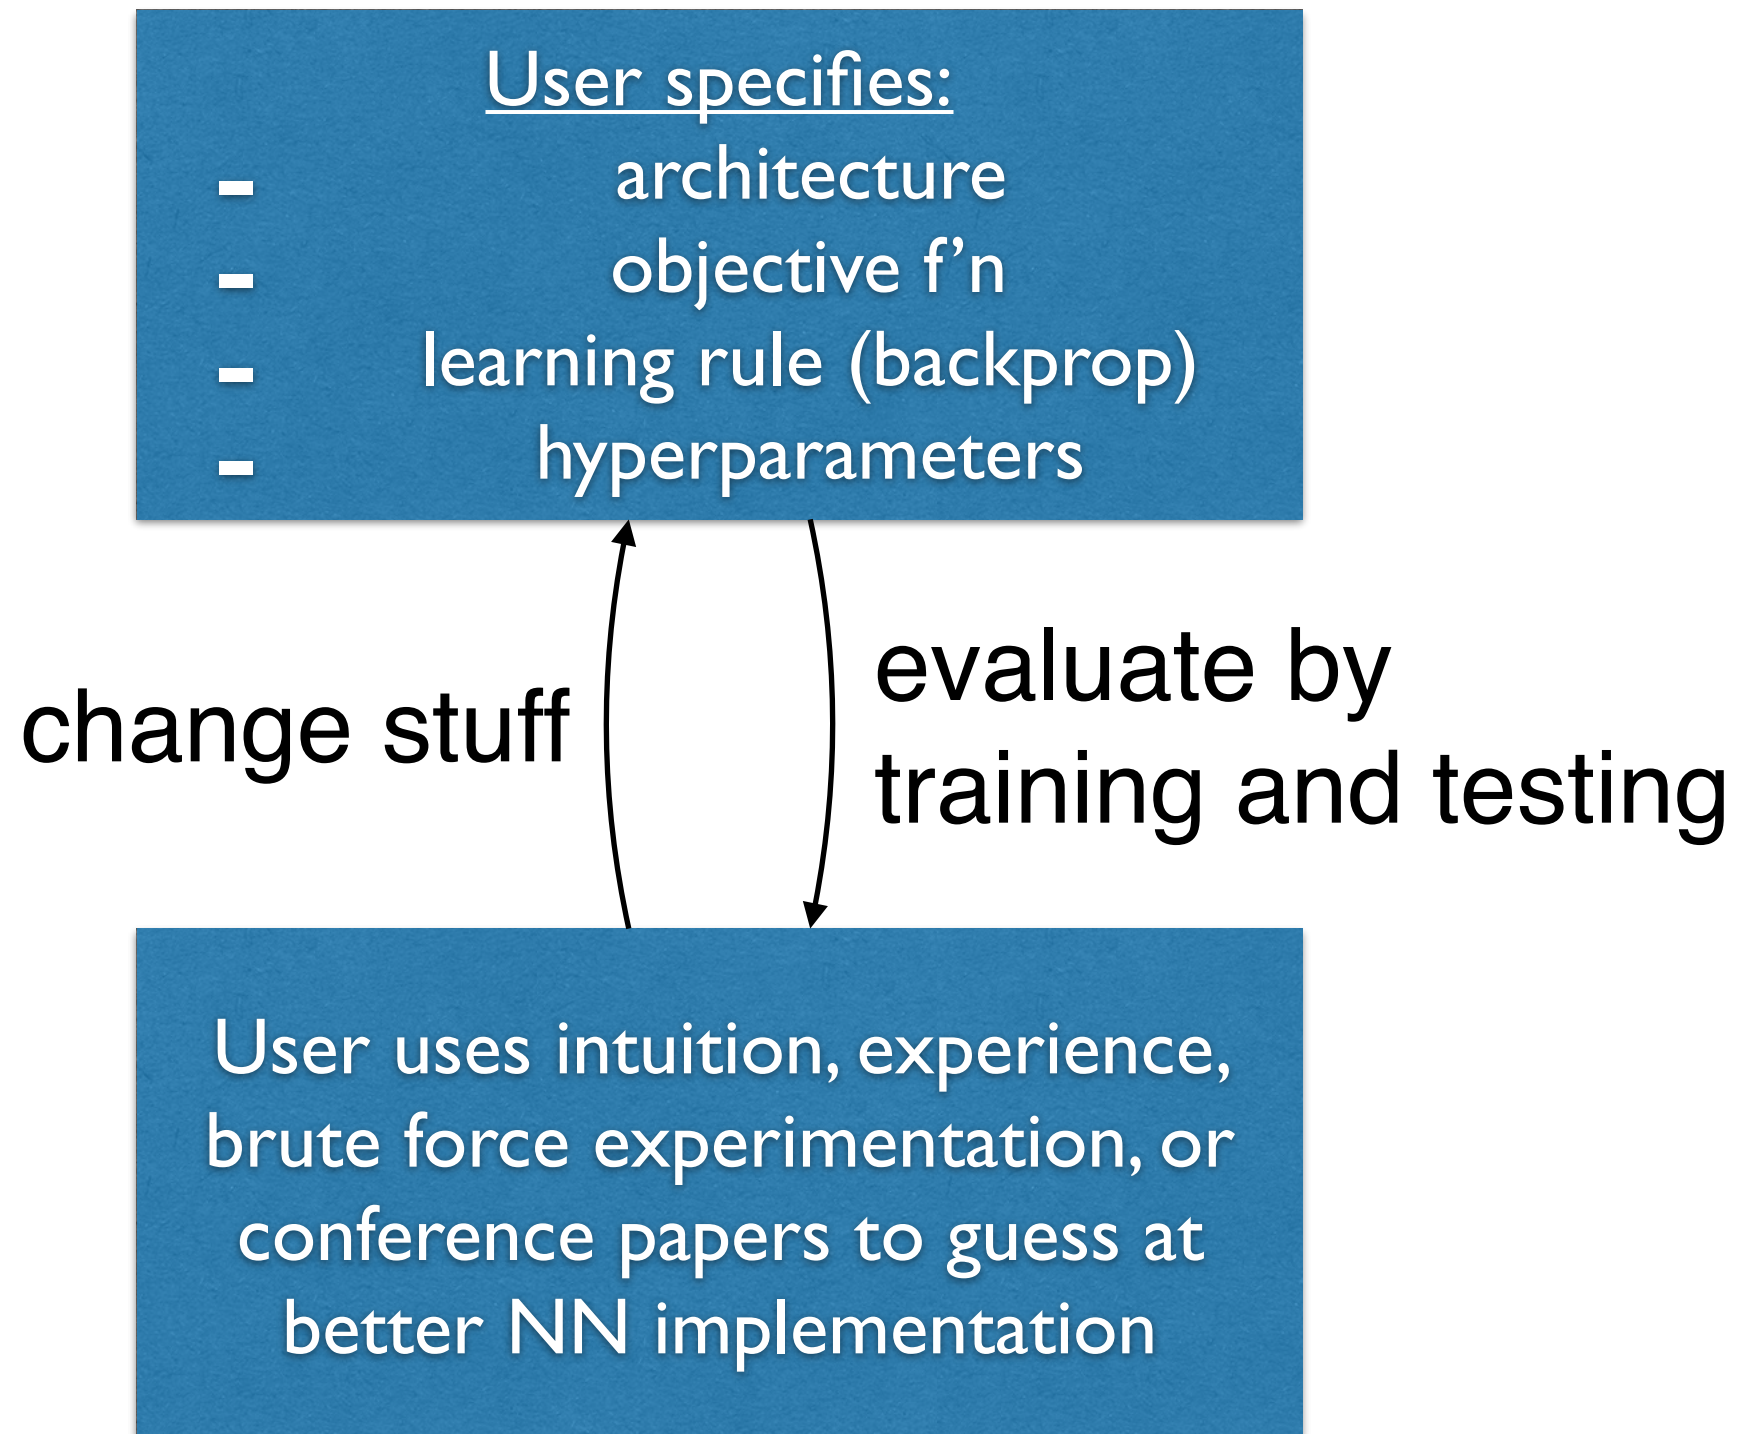
# **AutoML-Zero: Evolving Machine Learning Algorithms From Scratch**

Esteban Real<sup>\*1</sup> Chen Liang<sup>\*1</sup> David R. So<sup>1</sup> Quoc V. Le<sup>1</sup>

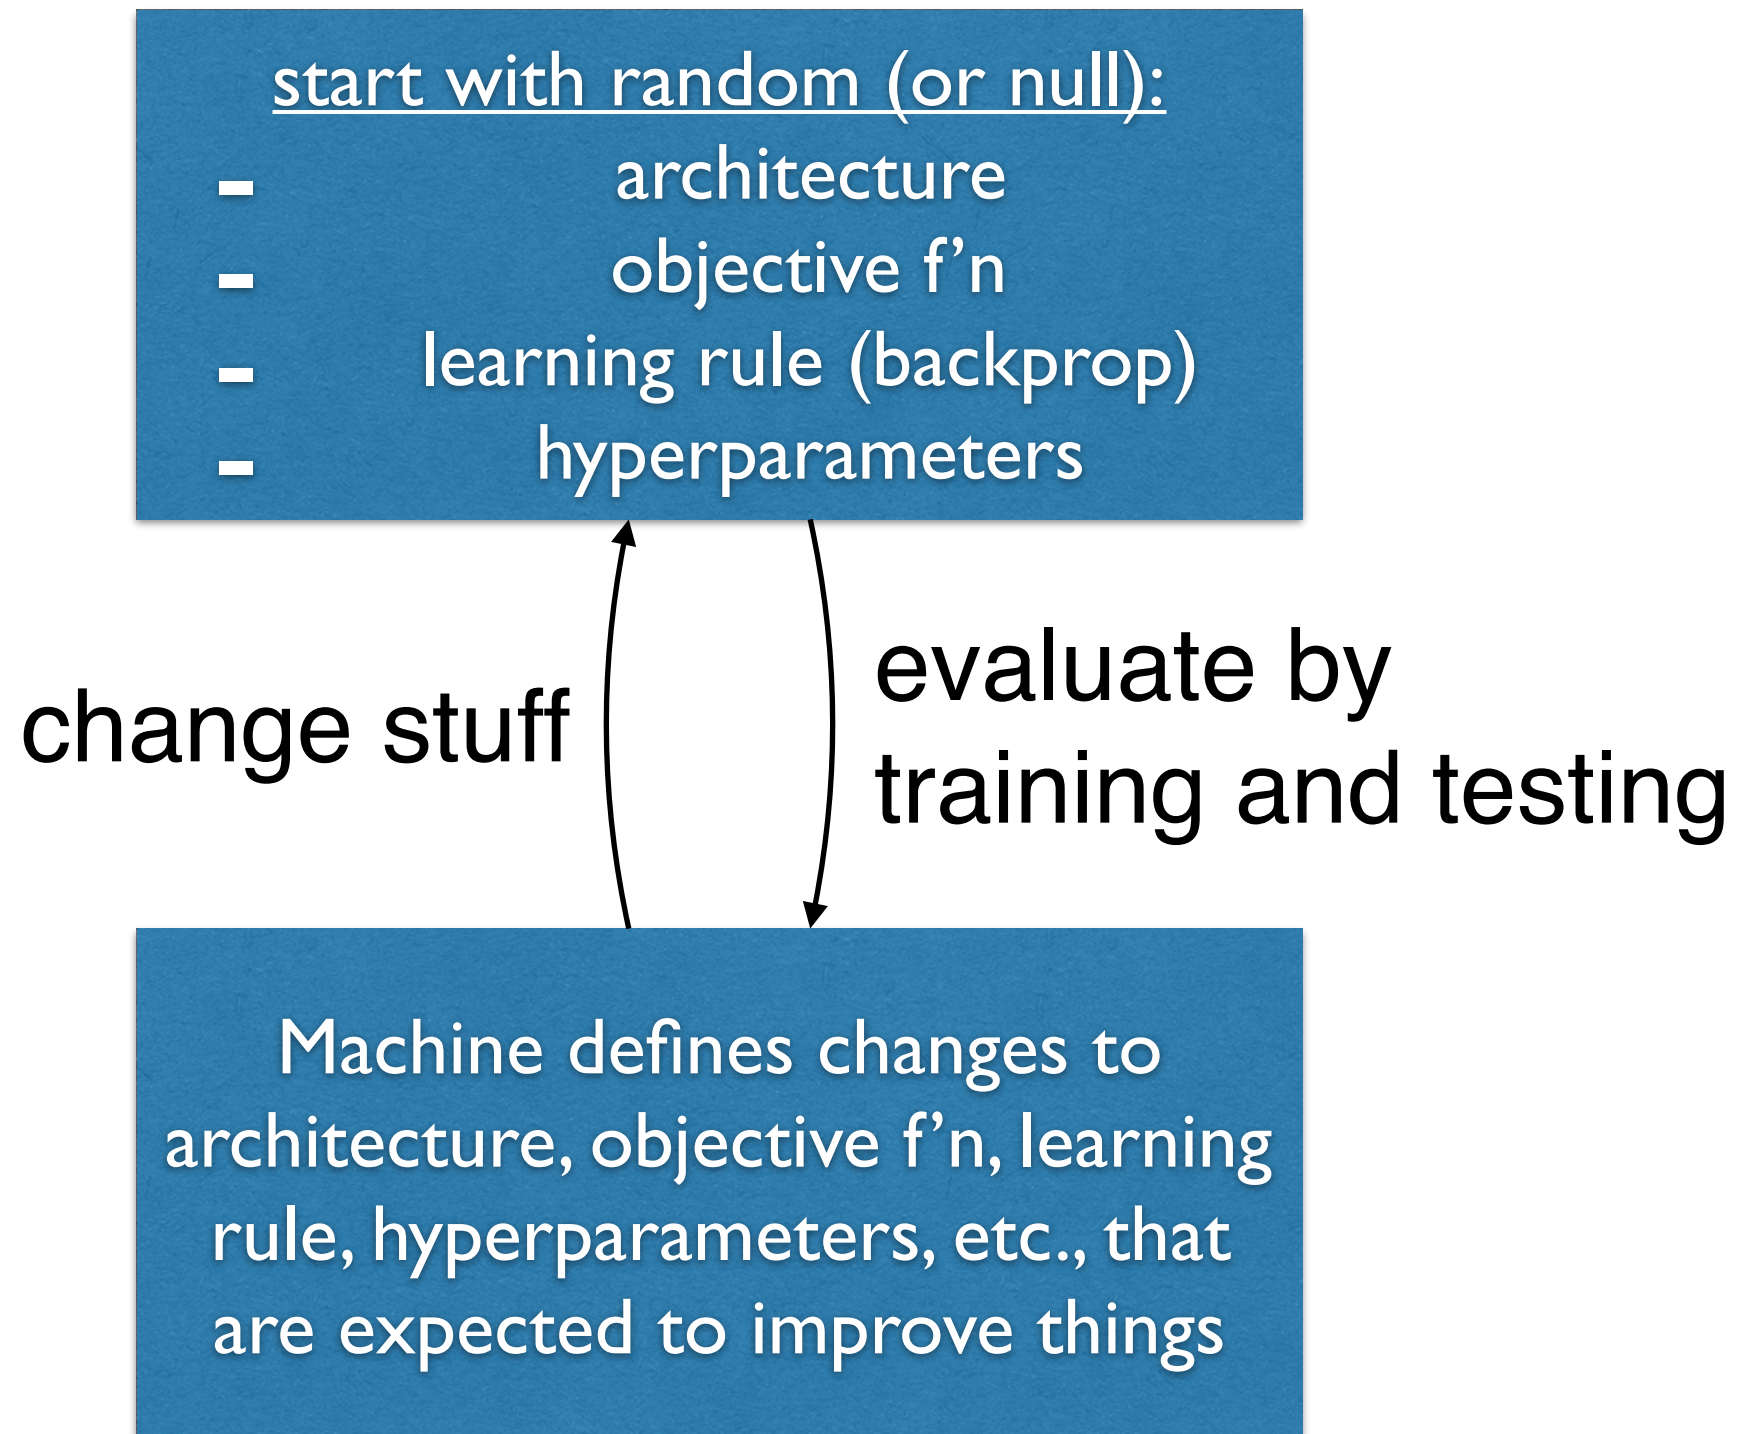
Presenter: Joel Zylberberg

[www.jzlab.org](http://www.jzlab.org)

# The “usual” Neural Nets Approach



# The “autoML” Approach



# Context on AI and Machine Learning

[VJ Manickam, 2020]

## THE MAIN CLASSIFICATION OF AI



### STRONG AI

AKA artificial general intelligence, an AI system with generalized human cognitive abilities. When presented with an unfamiliar task, it has enough intelligence to find a solution.

This doesn't exist yet  
(but is desirable and/or dangerous)



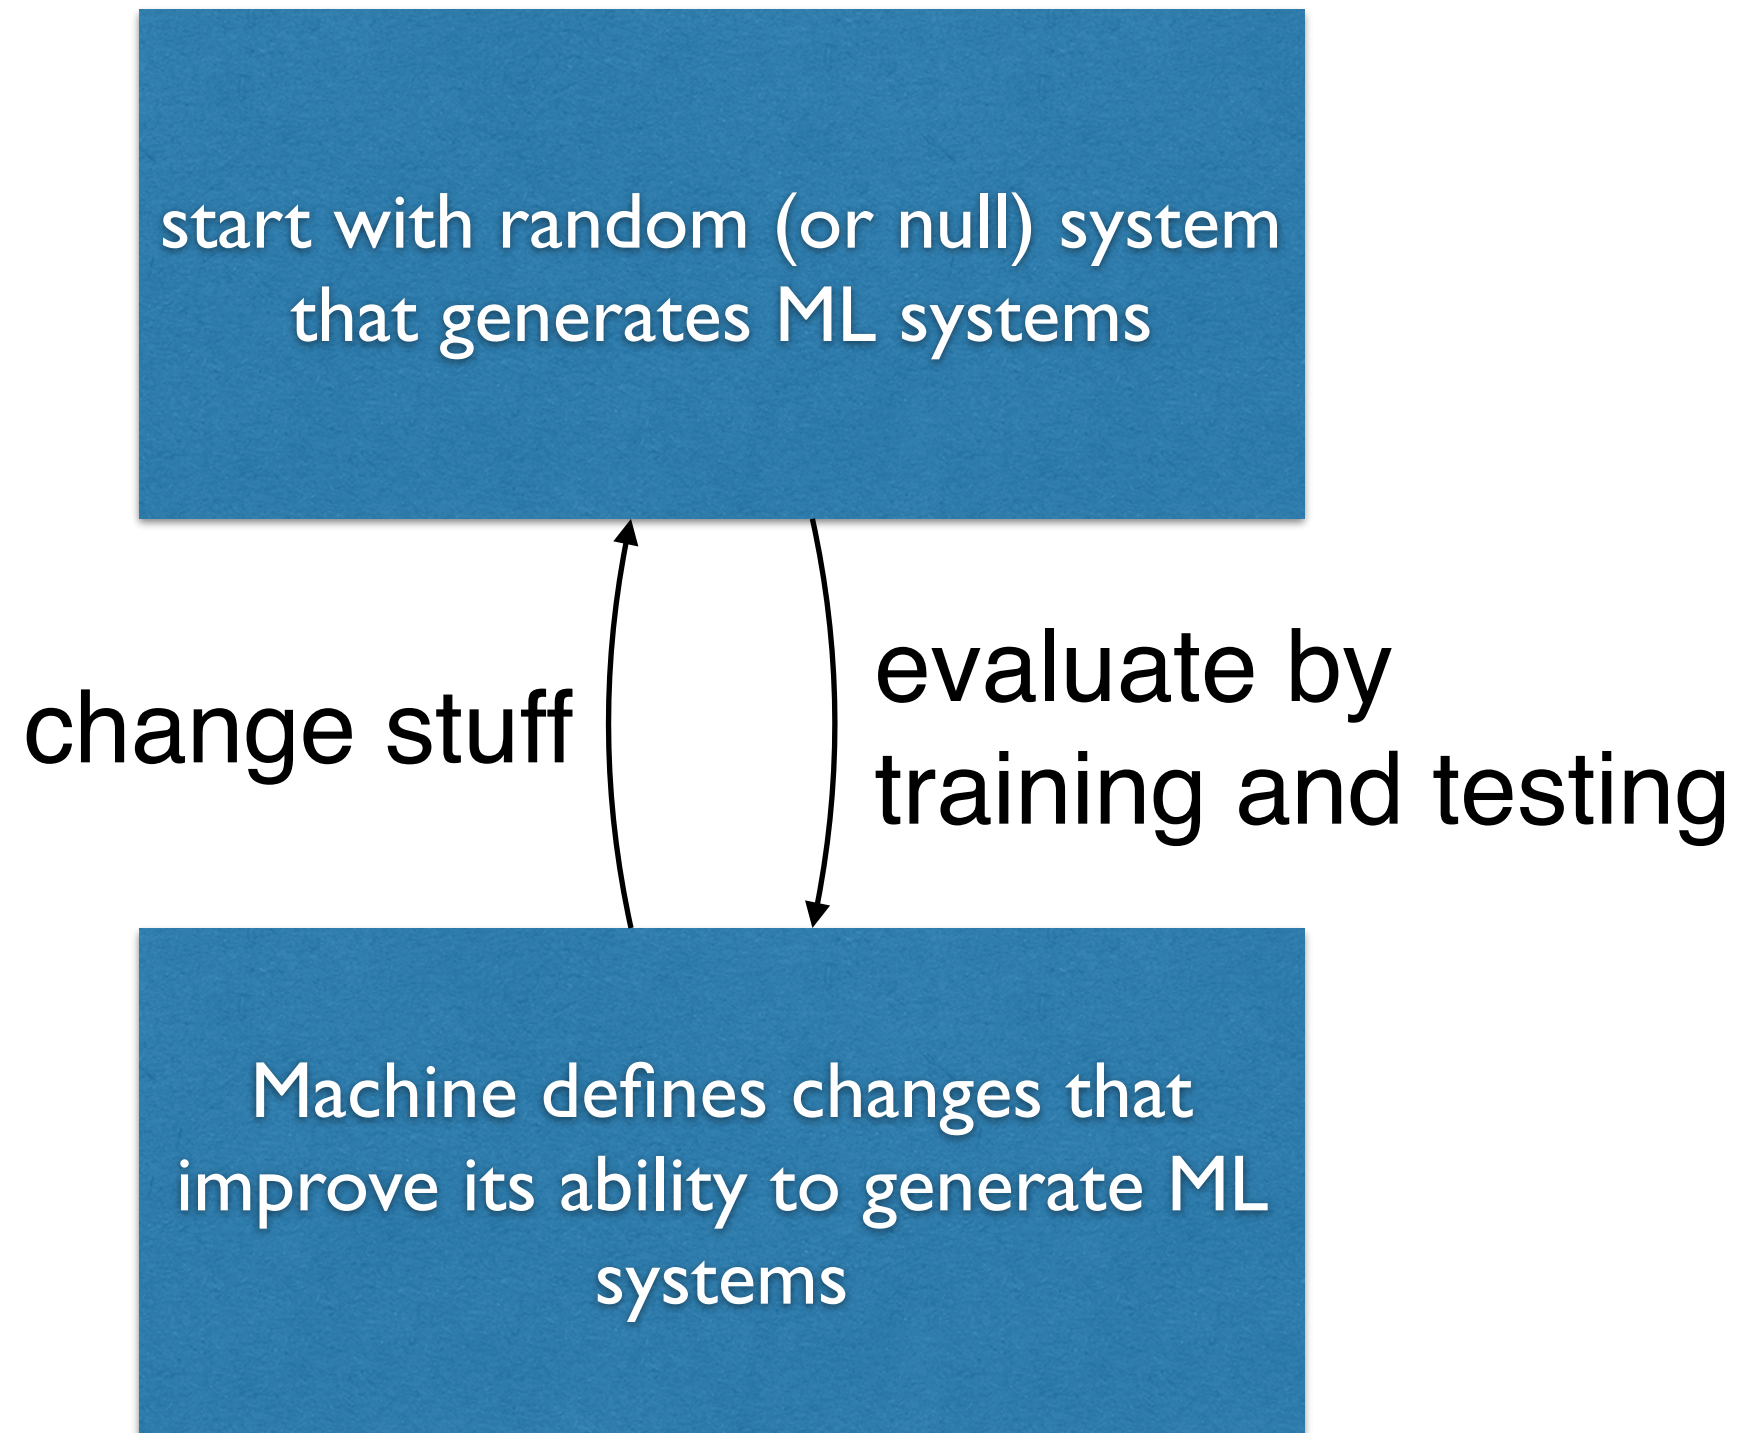
### WEAK AI

AKA narrow AI, an AI system that is designed and trained for a particular task. Example: a virtual personal assistant, such as Apple's Siri.

This is a product of hand-built NNs  
(and is rapidly improving)

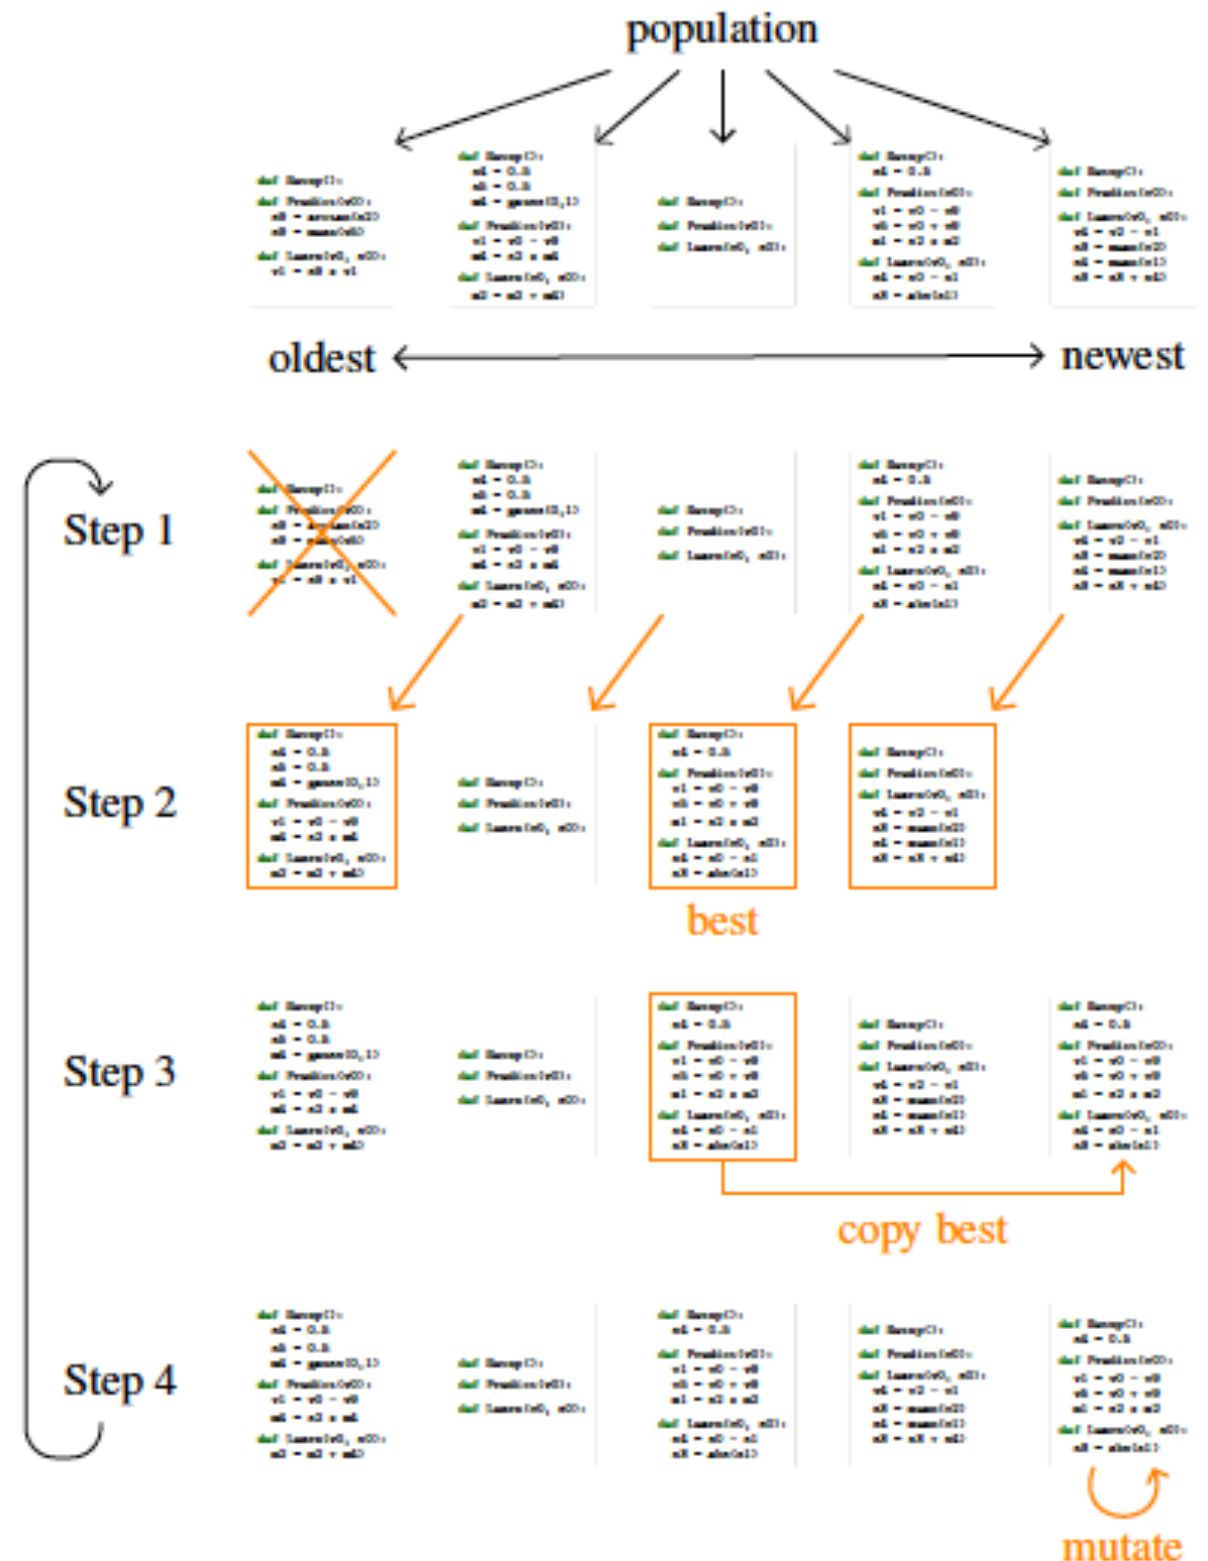


# The dreams of field

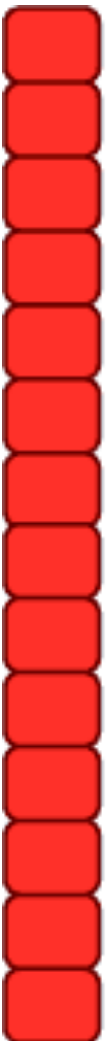


# Basic idea of the paper:

- 1) define an algorithmic procedure for writing ML algorithms
- 2) use simulated evolution to evolve a population of these ML algorithms (best ones propagated forward with mutation)
- 3) spend **tons** of compute resources and wait awhile



# Core pieces of AutoML-Zero programs

<u>Virtual Memory</u>	<u>Instructions</u>	<u>Component Functions</u>
 S1 S2 . .	<b>operation</b> and associated <b>memory address(es)</b>  E.g., Read input from <b>S1,S2</b> <b>multiply them together,</b> write answer to <b>S3</b>	setup, predict, learn  Each component function is a series of instructions

# Example programs (the big population of these evolves)

```
def Setup():  
    s4 = 0.5  
    s5 = 0.5  
    m4 = gauss(0,1)  
def Predict(v0):  
    v1 = v0 - v9  
    m4 = s2 * m4  
def Learn(v0, s0):  
    m2 = m2 + m4
```

```
def Setup():  
def Predict(v0):  
def Learn(v0, s0):
```

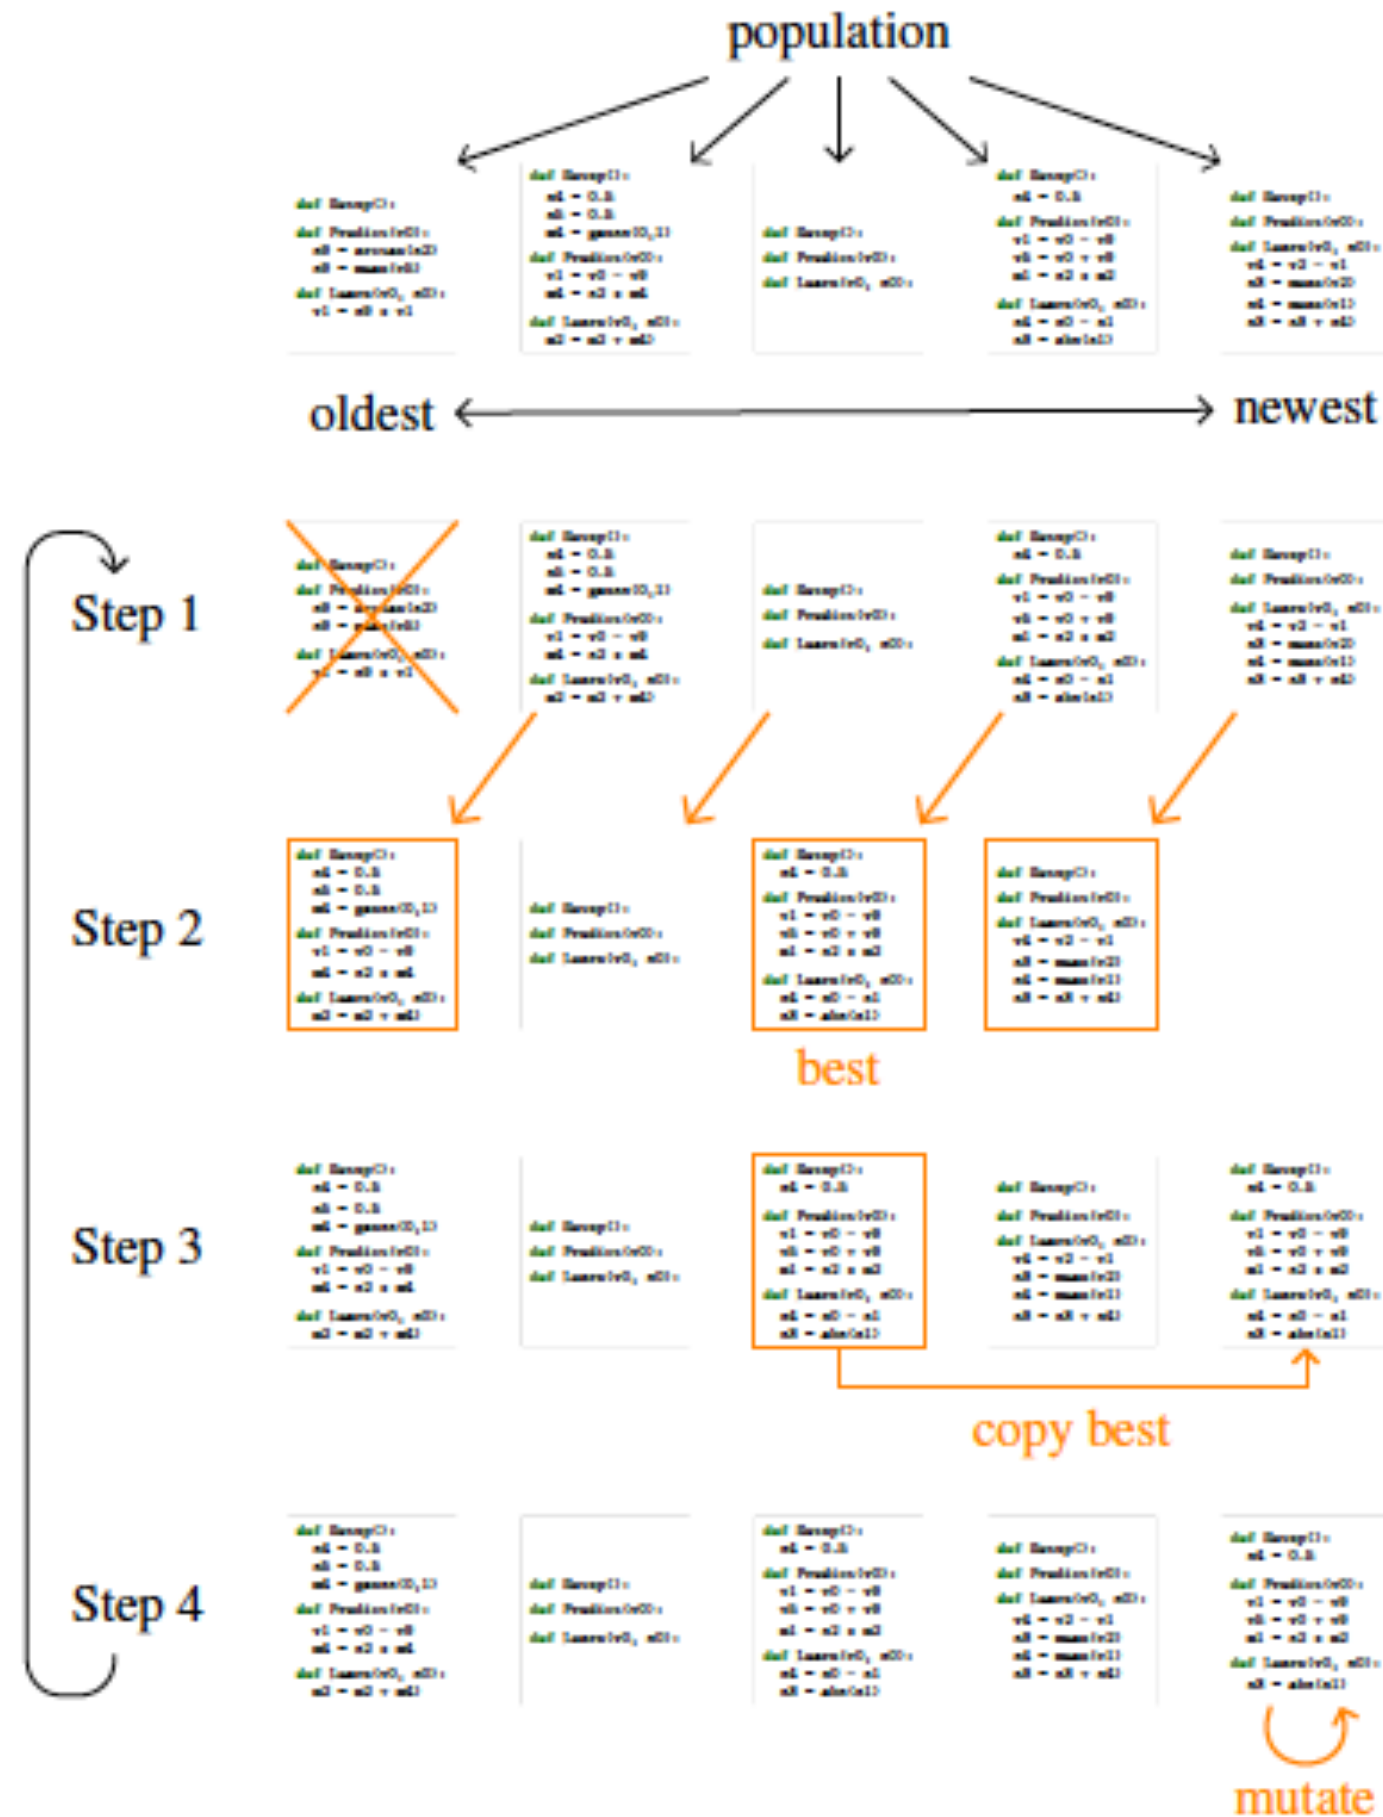
```
def Setup():  
    s4 = 0.5  
def Predict(v0):  
    v1 = v0 - v9  
    v5 = v0 + v9  
    m1 = s2 * m2  
def Learn(v0, s0):  
    s4 = s0 - s1  
    s3 = abs(s1)
```

```
def Setup():  
def Predict(v0):  
def Learn(v0, s0):  
    v4 = v2 - v1  
    s3 = mean(v2)  
    s4 = mean(v1)  
    s3 = s3 + s4
```

this one  
doesn't do  
anything



# Simulated evolution procedure



# Simulated evolution procedure hinges on evaluation (for finding the best model in each population)

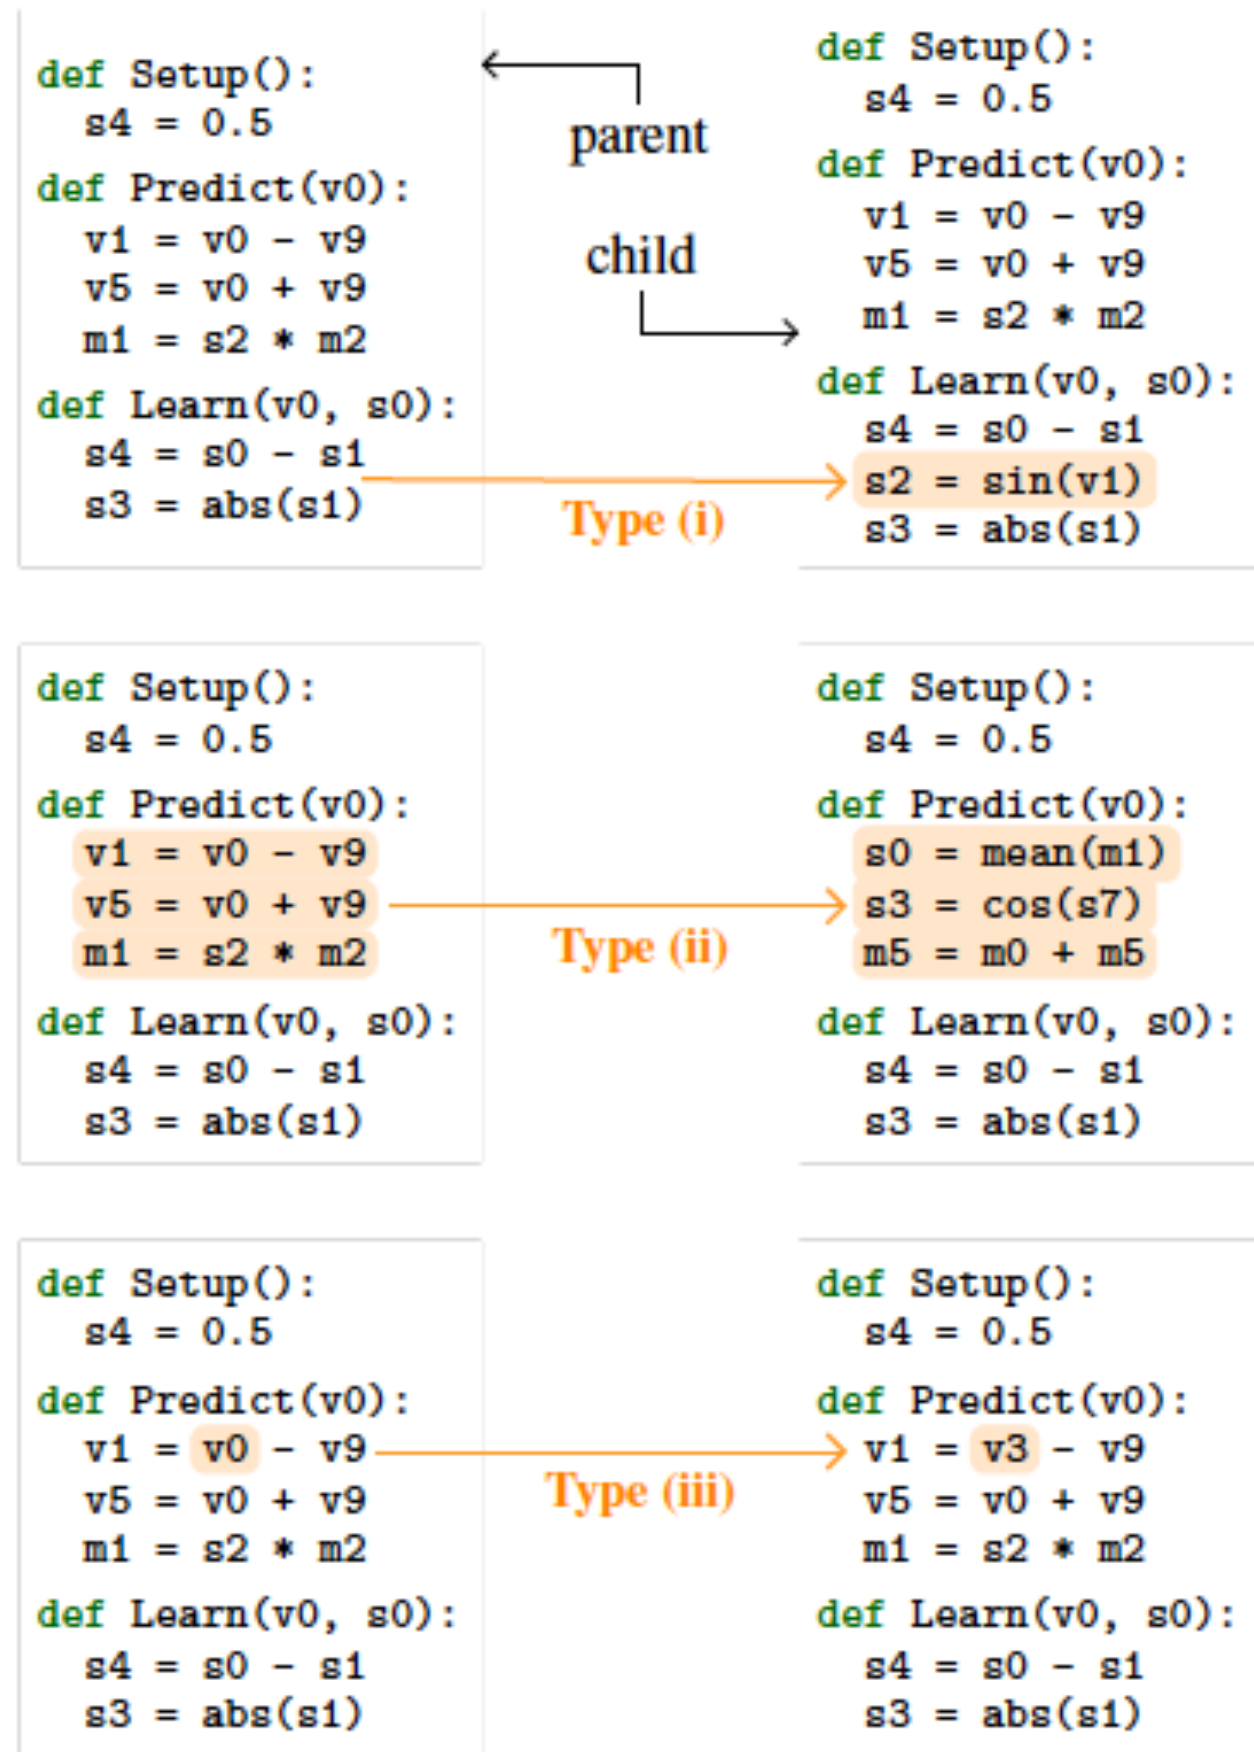
## Evaluation Procedure

```
# (Setup, Predict, Learn) is the input ML algorithm.
# Dtrain / Dvalid is the training / validation set.
# sX/vX/mX: scalar/vector/matrix var at address X.
def Evaluate(Setup, Predict, Learn, Dtrain, Dvalid):
    # Zero-initialize all the variables (sX/vX/mX).
    initialize_memory()
    Setup() # Execute setup instructions.
    for (x, y) in Dtrain:
        v0 = x # x will now be accessible to Predict.
        Predict() # Execute prediction instructions.
        # s1 will now be used as the prediction.
        s1 = Normalize(s1) # Normalize the prediction.
        s0 = y # y will now be accessible to Learn.
        Learn() # Execute learning instructions.

    sum_loss = 0.0
    for (x, y) in Dvalid:
        v0 = x
        Predict() # Only execute Predict(), not Learn().
        s1 = Normalize(s1)
        sum_loss += Loss(y, s1)

    mean_loss = sum_loss / len(Dvalid)
    # Use validation loss to evaluate the algorithm.
    return mean_loss
```

Figure 3: Mutation examples. Parent algorithm is on the left; child on the right. (i) Insert a random instruction (removal also possible). (ii) Randomize a component function. (iii) Modify an argument.



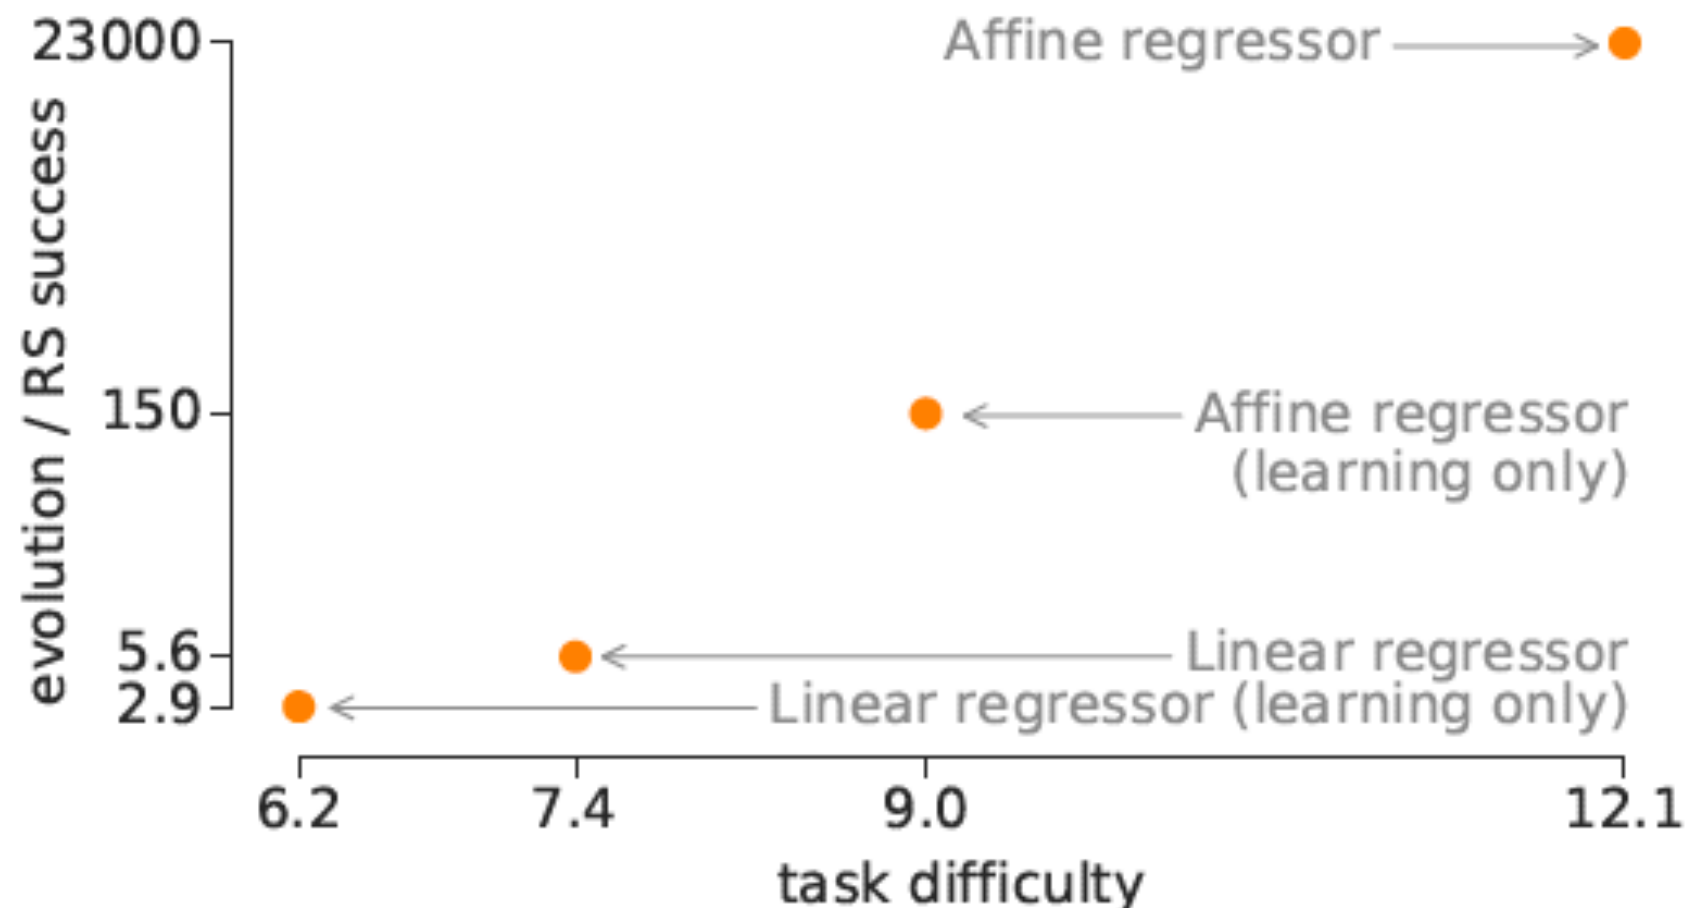
Key thing: the space of operations their systems can use / try is NOT limited to neural network operations. Includes a long list of “reasonable & simple” math operations

Now that the framework is in place...

Is it necessary? (Would Random Search — RS  
— do just as well?)

On simple regression problems, evolution approach finds more “good” solutions than RS. Algorithm search space is big so RS is bad.

Figure 4: Relative success rate of evolution and random search (RS). Each point represents a different task type and the x-axis measures its difficulty (defined in main text). As the task type becomes more difficult, evolution vastly outperforms RS, illustrating the complexity of AutoML-Zero when compared to more traditional AutoML spaces.





## Experiment 2:

- start with a random 2-layer neural net. Use it to generate input-output pairs
- evaluate programs in autoML based on whether they can duplicate that NNs outputs
- after evolution, read the best program and add comments to it

```
# sX/vX/mX = scalar/vector/matrix at address X.
# The C_ (eg C1) are constants tuned by search.
# "gaussian" produces Gaussian IID random numbers.
def Setup():
    # Initialize variables.
    m1 = gaussian(-1e-10, 9e-09) # 1st layer weights
    s3 = 4.1 # Set learning rate
    v4 = gaussian(-0.033, 0.01) # 2nd layer weights
def Predict(): # v0=features
    v6 = dot(m1, v0) # Apply 1st layer weights
    v7 = maximum(0, v6) # Apply ReLU
    s1 = dot(v7, v4) # Compute prediction
def Learn(): # s0=label
    v3 = heaviside(v6, 1.0) # ReLU gradient
    s1 = s0 - s1 # Compute error
    s2 = s1 * s3 # Scale by learning rate
    v2 = s2 * v3 # Approx. 2nd layer weight delta
    v3 = v2 * v4 # Gradient w.r.t. activations
    m0 = outer(v3, v0) # 1st layer weight delta
    m1 = m1 + m0 # Update 1st layer weights
    v4 = v2 + v4 # Update 2nd layer weights
```

Figure 5: Rediscovered neural network algorithm. It implements backpropagation by gradient descent. Comments added manually.



Experiment 3:

- evaluate programs in autoML based on whether they can label CIFAR-10 images. (Simplified version, tbh)
- at a few stages of evolution, read the best program and add comments to it

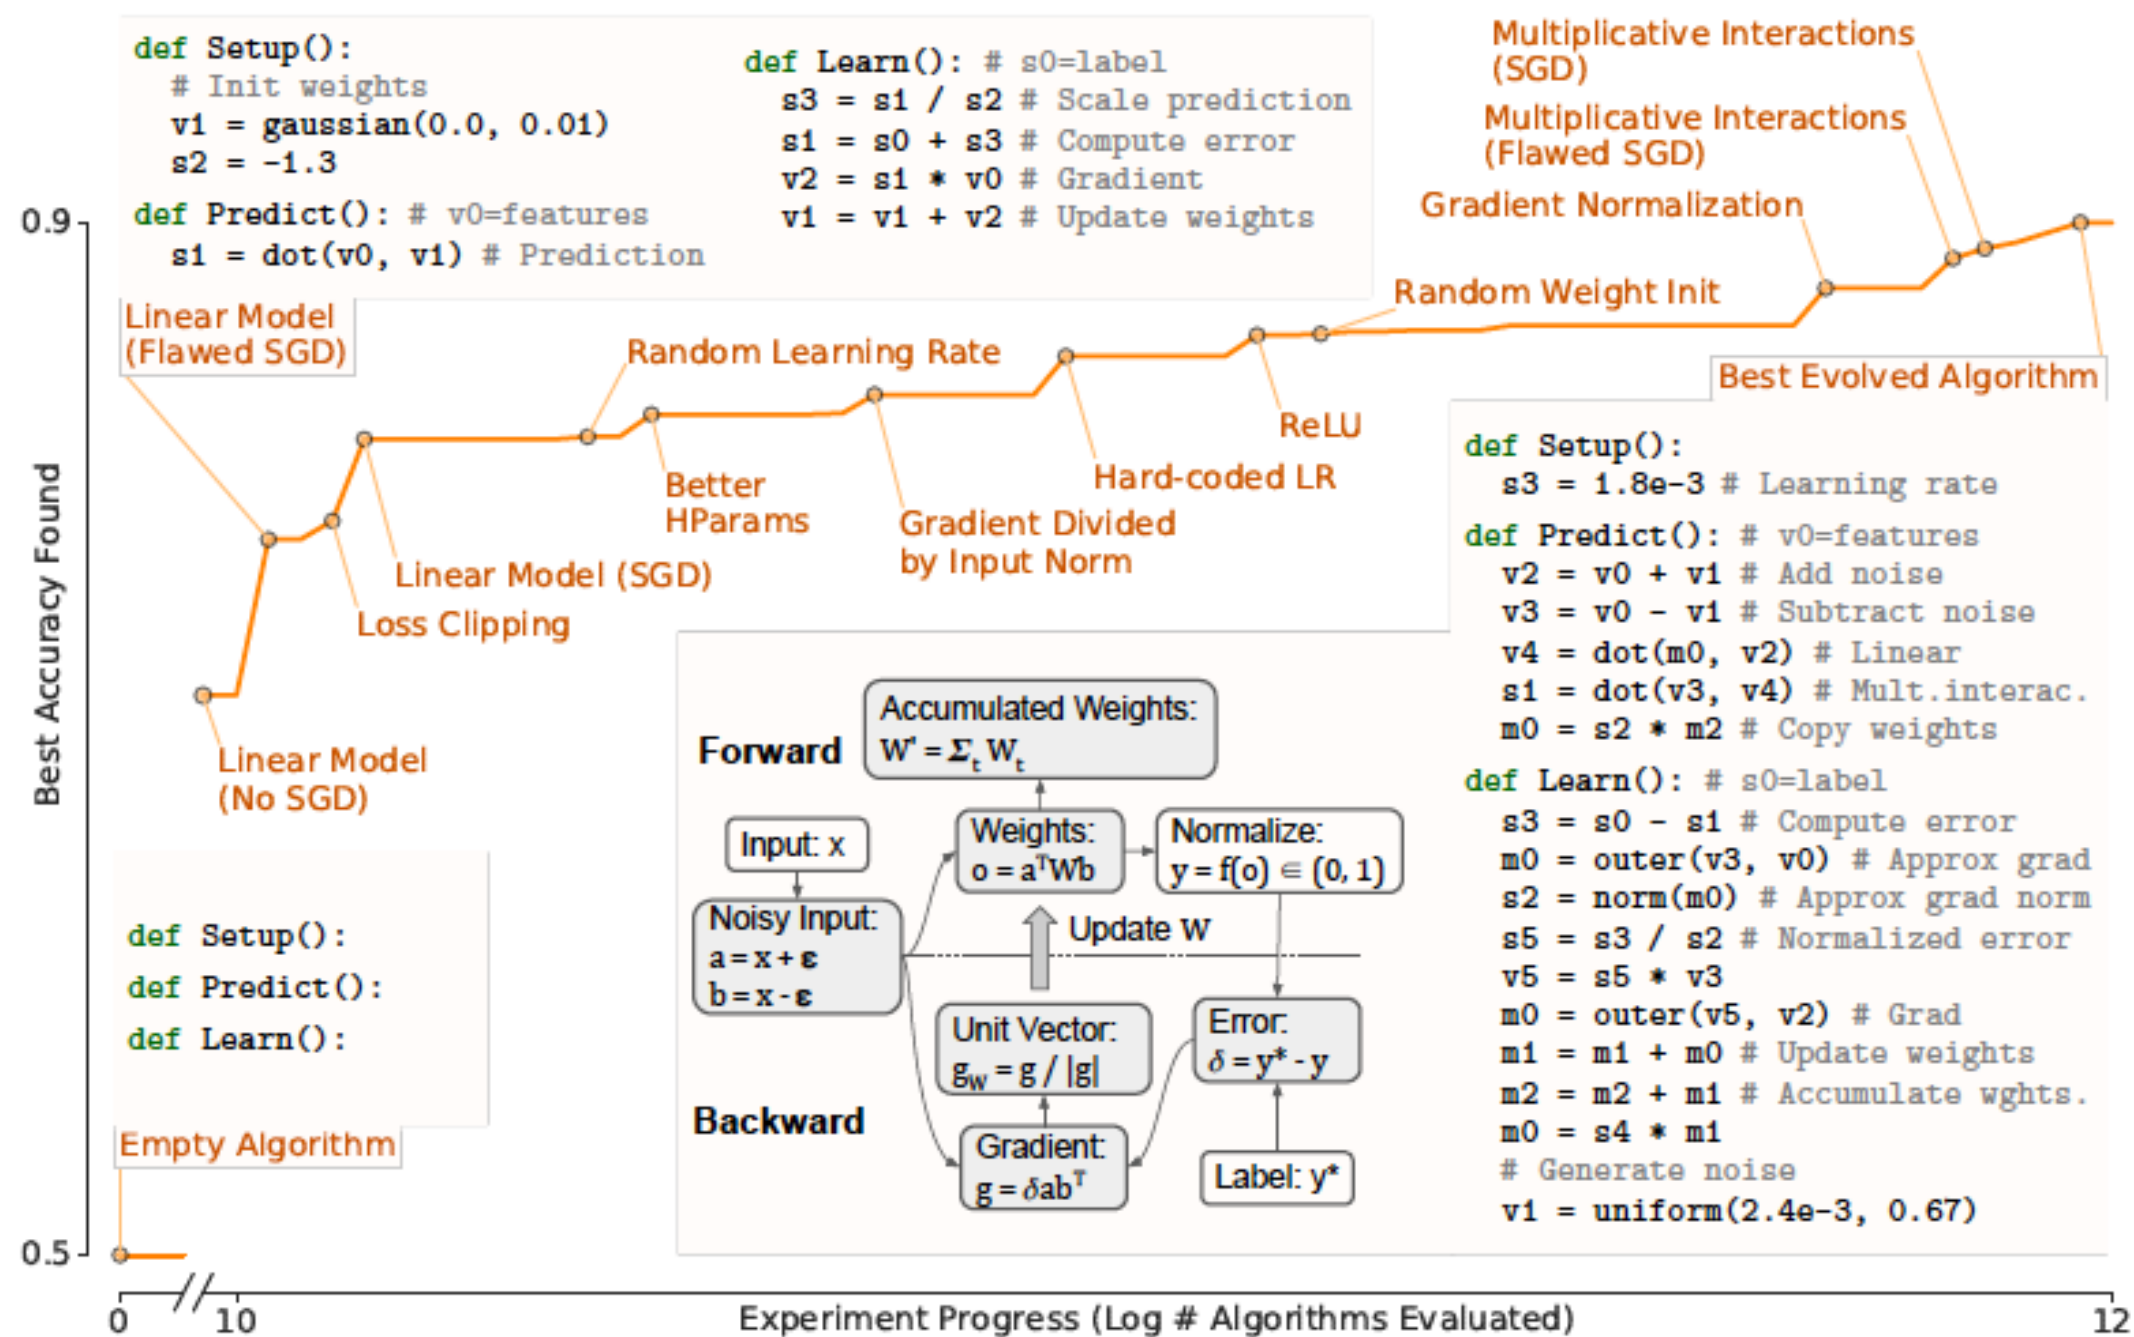


Figure 6: Progress of one evolution experiment on projected binary CIFAR-10. Callouts indicate some beneficial discoveries. We also print the code for the initial, an intermediate, and the final algorithm. The last is explained through the flow diagram. It outperforms a simple fully connected neural network on held-out test data and transfers to features 10x its size. Code notation is the same as in Figure 5.

There are a couple of other points in the paper, but (to me) not as important.

Discussion points:

- should we see autoML as a way to generate the actual ML algorithms we want, or instead look at the algos autoML finds, identify good innovations within them, and then use those innovations in hand-designed algos?
- why not implement a “smarter” optimization than evolution? (E.g., backprop or RL)
- what’s preventing the next stage of recursion: autoML that generates autoML that...
- should we be disappointed that the key innovations autoML-zero found are ones we already knew (e.g., is the “discovery” potential of the tool limited by our ability to interpret the programs it generates?)